

The Procurement Exception Commander


Disruption to corrective action.

Outcome: cost avoided & time-to-recovery

THE SCENARIO

Meet Wei, a supply planner. An alert lands: a supplier will miss a delivery, demand has outrun safety stock, or a shipment will miss a port cut-off. The alert is not the damage; the manual scramble that follows is, pulling data from three systems and chasing three teams to execute a fix.

HOW THE AI EMPLOYEE WORKS (EXAMPLE)



Assess impact and options in parallel, branch on severity, route high-impact actions to a human commander.

This is one example of how it could work, not a required design. Add or remove Operator Agents, change how each one works and the logic behind it, make the flow more optimal, and make sure it handles every kind of exception. Teams differ on the how; judges compare on the result.

ORCHESTRATOR AND OPERATORS (READ THIS)

Operator Agent

One AI Employee that does one job end to end (extract a record, screen for duplicates, draft an email). It takes an input, does the work, uses the integrations it needs, and returns a result. One Operator, one responsibility.

Orchestrator

An Operator that runs other Operators. It decides which to call and in what order (sequence, parallel, or branch), passes context between them, retries on failure, and escalates to a human when a decision needs one. The Orchestrator is the manager.

To call an Operator from an Orchestrator: inside the Orchestrator's workflow, name the Operator you want to run, for example "I want to trigger the Risk Screen Operator Agent" or "call the Extraction operator agent." Build each Operator on its own first, then build the Orchestrator that triggers them by name, branches, retries, and escalates. One giant workflow that does everything is a single mega-agent and does not pass.

WHAT YOU BUILD

An AI Employee that treats a disruption like an incident: it maps the blast radius, sources alternatives, weighs the trade-offs, and executes the fix, while a human commander approves anything high-impact before customers hear about it.

THE DATA YOU GET

A Coupa-style procurement export: suppliers, contracts, purchase order headers and lines, order confirmations, plus inventory and demand feeds and inbound disruption notices, as a workbook plus CSV tables.

WORKING WITH THE DATA

The dataset is how you exercise your integrations, not a file to read from. Either seed the systems you connect with it (load the records into your CRM, ticket tool, database, or an Airtable base), or park it in a store like SharePoint, OneDrive, Box, or Dropbox and point an integration at it, then pull the data back through that integration. Either way the data reaches your Operators through a live integration, at least three: one channel, one system of record, and a live exception to a person.

A way in

A disruption notice via email or Teams.

A system of record

Orders and inventory in Airtable or Supabase; tasks in Jira or Asana.

A human loop

A Slack or Teams approval to the commander before high-impact actions.

Note: Google integrations (Drive, Sheets, Gmail, Calendar) are in beta for this event, so use alternatives such as Airtable, Supabase, SharePoint, OneDrive, Box, Dropbox, and Outlook.

MANDATORY IN ROUND 1

1. Orchestrator and Operators. An Orchestrator coordinating at least two distinct Operators, showing parallel, branching, or stateful behavior. A single agent that does everything does not pass.

2. Live integrations. Work through real connected systems, not a file on your laptop. Minimum three integrations across two categories (one channel, one system of record), with a live exception to a human.

DO'S AND DON'TS

Do

- Build separate Operators for separate jobs, and one Orchestrator that coordinates them.
- Put the data into a real system, or pull it from SharePoint or OneDrive, and read and write through live integrations.
- Route genuine exceptions to a human at the Auto Workbench.
- Handle the messy data and the mismatches between records.
- Keep logic configurable without code.

Don't

- Build one giant single agent that does everything.
- Read the dataset file directly and process it with no integration.
- Hardcode to these sample rows; judging uses fresh data.
- Crash or invent a value on a missing field; pause and escalate.
- Skip the human loop or fake the demo.

HOW YOU'LL BE JUDGED

Four things carry the score. Business output weighs most; the rest show the build is real, adaptable, and demonstrable.

Business output

40

Did the AI Employee actually do the job on the data given, and is the result correct? A real, quantified result on the outcome metric (cost avoided & time-to-recovery), produced live, is the heart of the score.

Customizability

20

Can a business user quickly change the logic, thresholds, routing, or templates, or point the AI Employee at different incoming data, and still have it perform, without rewriting code?

Technical architecture

20

How the Operator Agents are orchestrated: the Orchestrator coordinates them, passes context cleanly, retries on failure, and escalates to a human when a decision needs one.

The demo

20

A clear, live walkthrough on the Auto Manager Console, run end to end on data the team did not pre-arrange.

Bonus. Extra useful downstream actions, more Operators, and genuine innovativeness and creativity beyond the brief score above a baseline build. Solve the stated problem first, then extend it.

Built to be tested, not staged. During judging the data changes: each build is run live on a separate, unseen dataset, and judges may test with several. Your AI Employee should be universal, able to ingest whatever data arrives, understand it, and handle the exceptions, rather than tuned to these sample rows.

FAQ

Do I need deep domain expertise?

No. The scenario and dataset give you what you need; focus on the workflow and the exceptions, not the domain.

Can a single agent solve this?

No. The track needs an Orchestrator coordinating multiple Operator Agents, with at least one live exception routed to a human at the Auto Workbench.

How is my build judged?

On a quantified result on the outcome metric, produced live on a separate dataset your agent has not seen. Solve the general problem, not these specific rows.

Not apps. Not agents. AI Employees run the Autos.